

Predicting Stackoverflow tags

02807 Final project

November 29, 2016

1 Introduction

In this project a model predicting tags associating a given text will be constructed. The problem will be handled as a classification problem, hence a classifier will be trained.

The classifier will be trained on posts from Stackoverflow where associated tags have been given.

Initially only the top 20 frequent tags and posts containing these tags will be used for training and prediction.

Finally the model will be evaluated using all tags and posts.

During implementation and debugging a subset of the data will be used.

2 The data

The dataset consists of two XML files, one containing all possible tags and their corresponding counts, and one containing posts with *title*, *body*, *tags* and some meta data.

The total size of the files are approximately 49GB in uncompressed format.

3 Methods

In this section the different methods and steps in the process will be explained.

3.1 Preprocessing

The preprocessing step regards the transforming of questions in an XML file to processed questions in a .csv file. This also includes disregarding questions that does not have any of the top N tags attached. At the same time two other .csv

files are created: One containing all unique words in the extracted questions, and one containing the unique tags used.

The processing of each questions contains the following steps (code can be found in appendix A.1)

1. Replace all links with `<link>`
2. Remove certain unwanted symbols
3. Remove suffix from words (e.g. *haven't* \rightarrow *have*)
4. Remove line breaks
5. Replace digits with `<digit>`
6. Remove double whitespaces
7. Reduce words to their word stem (e.g. *lines* \rightarrow *line*)
8. Lemmatize words (e.g. *better* \rightarrow *good*)

Finally the unique words used as word dictionary were filtered by removing words occurring in more than 50% of the questions and words occurring in less than 0.1% of the questions. Also english stop words were removed making use of the NLTK library.

All these steps are used in order to reduce the dimensionality of the word space without really removing much information. Here the assumption is, that e.g. words like *better* and *good* kind of adds the same meaning to the sentence, and the same with e.g. two numbers.

3.2 Distributed file loading

Since the size of the final processed questions file is approximately 11GB, it will not be feasible to load into memory on most laptops. Therefore it will be necessary to load the file in smaller chunks.

The following code illustrates how the file `posts.csv` can be divided into byte-chunks. I.e. the following generator yields a list of tuples (`from_byte`, `size`) where `from_byte` is the index in the file in bytes and `size` is the size of the given chunk in bytes.

```

1 | with open('posts.csv', 'rb') as f:
2 |     while True:
3 |         start = f.tell()
4 |         f.seek(chunk_size, 1)
5 |         s = f.readline()
6 |         yield start, f.tell() - start
7 |         if not s: break

```

The `chunk_size` is a given minimum size of each chunk. The `f.readline()` makes sure the chunk ends at the end of a line.

A chunk of lines from the file can then be loaded using the following lines

```

1 | # Seek to chunk start bytes
2 | f.seek(from_bytes)
3 |
4 | # Read end of chunk until end of line
5 | chunk = f.read(size)
6 |
7 | # Split in lines (Removing the last newline)
8 | lines = chunk.rstrip('\n').split('\n')

```

3.3 K-means clustering

3.3.1 Serial

The regular serial in-memory version of K-means clustering algorithm is shown in algorithm 1.

Algorithm 1 Serial K-means clustering algorithm

```

1: procedure KMEANSCLUSTERING( $X, K$ )
2:   # Initialize cluster centers
3:   for  $k = 0$  to  $K - 1$  do
4:      $\mu_k \leftarrow$  random point in  $X$ 
5:   # Run iterations
6:   while  $iter < max\_iter$  do
7:     # Update cluster means
8:      $\mu_{old} = \mu$ 
9:     for  $k = 0$  to  $K - 1$  do
10:       $C_k \leftarrow$  {Points in  $X$  closest to  $\mu_k$ }
11:       $\mu_k \leftarrow \frac{1}{|C_k|} \sum_{x_i \in C_k} x_i$ 
12:    # Check convergence criteria
13:     $norm \leftarrow \|\mu - \mu_{old}\|$ 
14:    if  $norm < \epsilon$  then
15:      break

```

3.3.2 Distributed

The proposed distributed K-means clustering algorithm, which loads the data matrix X in chunks, is shown in algorithm 2.

Algorithm 2 Distributed K-means clustering algorithm

```

1: procedure KMEANSCLUSTERINGDISTRIBUTED( $X, K$ )
2:    $\#$  Initialize cluster centers
3:   for  $k = 0$  to  $K - 1$  do
4:      $\mu_k \leftarrow$  random point in  $X$ 
5:    $\#$  Run iterations
6:   while  $iter < max\_iter$  do
7:      $\#$  Initialize shared cluster sums and cluster point counts.
8:     for  $k = 0$  to  $K - 1$  do
9:        $Csum_k \leftarrow 0$ 
10:       $Ccount_k \leftarrow 0$ 
11:      $\#$  Process each chunk in a distributed manner
12:     for all Chunks  $X_{chunk}$  in  $X$  do
13:       for  $k = 0$  to  $K - 1$  do
14:          $C_k \leftarrow \{\text{Points in } X_{chunk} \text{ closest to } \mu_k\}$ 
15:          $Csum_k \leftarrow Csum_k + \sum_{x_i \in C_k} x_i$ 
16:          $Ccount_k \leftarrow Ccount_k + |C_k|$ 
17:      $\#$  Gather results and update cluster means
18:      $\mu_{old} = \mu$ 
19:     for  $k = 0$  to  $K - 1$  do
20:        $\mu_k \leftarrow \frac{Csum_k}{Ccount_k}$ 
21:      $\#$  Check convergence criteria
22:      $norm \leftarrow \|\mu - \mu_{old}\|$ 
23:     if  $norm < \epsilon$  then
24:       break

```

3.3.3 Implementation

Simplified implementation of distributed K-means (only a single iteration is shown) see full code in appendix A.2:

```

1 | cluster_sums = {k: np.zeros((1, word_count)) for k in
   | range(0, K)}
2 | cluster_counts = {k: 0 for k in range(0, K)}
3 |
4 | for chunk in chunks:
5 |
6 |      $\#$  Load chunk lines to sparse matrix
7 |      $X = \text{chunk\_to\_sparse\_mat}(\text{chunk})$ 
8 |
9 |      $\#$  Get closest cluster indices
10 |    max_idx = sparse_matrix_to_cluster_indices( $X, \mu$ )
11 |
12 |     $\#$  Assign points to clusters
13 |    mu_subs = collections.defaultdict(list)
14 |    for i, k in enumerate(max_idx):
15 |        mu_subs[k].append( $X[i].\text{toarray}()$ )
16 |
17 |     $\#$  Compute sub-means
18 |    for k in range(0, K):

```

```
19     mu_sub = mu_subs[k]
20     if len(mu_sub) == 0:         continue
21     cluster_sums[k] += np.asarray(mu_sub).mean(axis=0)
22     cluster_counts[k] += 1
23
24 # Save old means
25 mu_old = np.array(mu, copy=True)
26
27 # Update means
28 for k in range(0, K):
29     count = cluster_counts[k]
30     if count == 0: continue
31     mu[k] = cluster_sums[k] / cluster_counts[k]
32
33 # Check convergence criteria
34 mu_norm = np.linalg.norm(mu - mu_old)
35
36 if mu_norm < epsilon:
37     print('Converged after %d iterations' % (iteration+1))
38     break
```

3.4 Parallel processing

4 Results

5 Discussion

6 Conclusion

A Code snippets

A.1 Preprocess text

```
1 import re
2 import Stemmer
3 from nltk.stem import WordNetLemmatizer
4 lemmatizer = WordNetLemmatizer()
5
6 # Precompile regular expressions
7 reg_links = re.compile(r'http[s]?://(?:[a-zA-Z]|[0-9]|[$-_@
8 .&+]|[!*\(\)\,\:]|(?:%[0-9a-fA-F][0-9a-fA-F]))+')
9 reg_digits = re.compile(r'\b\d+\b')
10 reg_spaces = re.compile(r'\s{2,}')
11
12 reg_symbols = re.compile(r'^A-Za-z0-9(),!?\'\`')
13 reg_symb_1 = re.compile(r',')
14 reg_symb_2 = re.compile(r'!')
15 reg_symb_3 = re.compile(r'\(')
16 reg_symb_4 = re.compile(r'\)')
17 reg_symb_5 = re.compile(r'\?')
18 reg_symb_6 = re.compile(r'\'')
19
20 reg_suf_1 = re.compile(r'\s')
21 reg_suf_2 = re.compile(r'\ve')
22 reg_suf_3 = re.compile(r'n\t')
23 reg_suf_4 = re.compile(r'\re')
24 reg_suf_5 = re.compile(r'\d')
25 reg_suf_6 = re.compile(r'\ll')
26
27 stemmer = Stemmer.Stemmer('english')
28 word_to_stem = {}
29 def stem_word(word):
30     if not word in word_to_stem:
31         word_to_stem[word] = stemmer.stemWord(word)
32     return word_to_stem[word]
33
34 word_to_lemma = {}
35 def lemmatize_word(word):
36     if not word in word_to_lemma:
37         word_to_lemma[word] = lemmatizer.lemmatize(word)
38     return word_to_lemma[word]
39
40 def clean_string(text):
41     # Replace links with link identifier
42     text = reg_links.sub('<link>', text)
43
44     # Remove certain symbols
45     text = reg_symbols.sub(' ', text)
46
47     # Remove suffix from words
48     text = reg_suf_1.sub(' ', text)
49     text = reg_suf_2.sub(' ', text)
```

```
49 text = reg_suf_3.sub(' ', text)
50 text = reg_suf_4.sub(' ', text)
51 text = reg_suf_5.sub(' ', text)
52 text = reg_suf_6.sub(' ', text)
53
54 # Remove "\"" from string
55 text = reg_symb_6.sub(' ', text)
56
57 # Replace breaks with spaces
58 text = text.replace('<br />', ' ')
59 text = text.replace('\r\n', ' ')
60 text = text.replace('\r', ' ')
61 text = text.replace('\n', ' ')
62
63 # Pad symbols with spaces on both sides
64 text = reg_symb_1.sub(' ', ' ', text)
65 text = reg_symb_2.sub(' ! ', text)
66 text = reg_symb_3.sub(' ( ', text)
67 text = reg_symb_4.sub(' ) ', text)
68 text = reg_symb_5.sub(' ? ', text)
69
70 # Replace digits with 'DIGIT'
71 text = re_digits.sub('<DIGIT>', text)
72
73 # Remove double whitespaces
74 text = re_spaces.sub(' ', text)
75 text = text.strip()
76
77 # Convert to lowercase
78 text = text.lower()
79
80 # Stem each word
81 text = ' '.join(stem_word(word) for word in text.split(' '))
82
83 # Lemmatize each word
84 text = ' '.join(lemmatize_word(word) for word in text.
    split(' '))
```

A.2 Distributed K-means

```
1 import math
2 import collections
3 import numpy as np
4 import multiprocessing
5 import time
6
7 import helpers
8 import config
9
10 # Read tags
11 tags, tag2idx, tag_count = helpers.read_tags()
12
13 # Read words
```

```
14 words, word2idx, word_count = helpers.read_words()
15
16 # Clusters
17 K = tag_count
18
19 # Initialize cluster centers
20 mu = np.random.rand(K, word_count)
21
22 # Get chunks
23 chunk_reader = helpers.ChunkReader(post_filename=config.
    paths.TRAIN_DATA_IDX, chunk_size=config.data.CHUNK_SIZE)
    # TODO: Change
24 chunks = [chunk for chunk in chunk_reader]
25 chunk_count = len(chunks)
26
27 # Split chunks across processes
28 n = math.ceil(chunk_count / config.algorithm.PROCESS_COUNT)
29 chunks_split = []
30 for i in range(0, len(chunks), n):
31     chunks_split.append(chunks[i:i+n])
32
33 # Initialize shared variable manager
34 manager = multiprocessing.Manager()
35 lock = multiprocessing.Lock()
36
37 # Define function to run in parallel
38 def process_chunks(chunks, word_count, K, mu, cluster_sums,
    cluster_counts, lock):
39     for chunk in chunks:
40
41         # Convert to sparse matrix
42         X, _ = helpers.chunk_to_sparse_mat(chunk, word_count)
43
44         if X is None: continue
45
46         # Get closest cluster indices
47         max_idx = helpers.sparse_matrix_to_cluster_indices(X, mu
    )
48
49         mu_subs = collections.defaultdict(list)
50         for i, k in enumerate(max_idx):
51             mu_subs[k].append(X[i].toarray())
52
53         # Compute sub-means
54         for k in range(0, K):
55             mu_sub = mu_subs[k]
56             if len(mu_sub) == 0: continue
57
58             with lock:
59                 cluster_sums[k] = cluster_sums[k] + np.asarray(
    mu_sub, dtype=np.float32).mean(axis=0)
60                 cluster_counts[k] += 1
61
62
```



```
63 for iteration in range(0, config.algorithm.MAX_ITER):
64     start = time.time()
65
66     cluster_sums = manager.dict({k: np.zeros((1, word_count))
67                                 for k in range(0, K)})
68     cluster_counts = manager.dict({k: 0 for k in range(0, K)})
69
70     # Init processes
71     processes = []
72     for i, chunk_list in enumerate(chunks_split):
73         p = multiprocessing.Process(target=process_chunks,
74                                   kwargs={
75                                       'chunks': chunk_list,
76                                       'word_count': word_count,
77                                       'K': K,
78                                       'mu': mu,
79                                       'cluster_sums': cluster_sums,
80                                       'cluster_counts': cluster_counts,
81                                       'lock': lock
82                                   })
83         processes.append(p)
84
85     # Start processes
86     for p in processes:
87         p.start()
88
89     #print('Started %d processes' % (len(processes)))
90
91     # Wait for processes to finish
92     for p in processes:
93         p.join()
94
95     # Save old means
96     mu_old = np.array(mu, copy=True)
97
98     # Update means
99     for k in range(0, K):
100         count = cluster_counts[k]
101         if count == 0: continue
102         mu[k] = cluster_sums[k] / cluster_counts[k]
103
104     # Check convergence criteria
105     mu_norm = np.linalg.norm(mu - mu_old)
106
107     print('Iteration %d took: %.4fs' % (iteration + 1, time.
108                                         time() - start))
109
110     if mu_norm < config.algorithm.EPSILON:
111         print('Converged after %d iterations' % (iteration+1))
112         break
113
114     # Determine cluster tags
115     cluster_tag_counts = {k: {tag: 0 for tag in range(0, K)} for
```

```
    k in range(0, K)}
114 for chunk in chunks:
115
116     # Convert to sparse matrix
117     X, tags = helpers.chunk_to_sparse_mat(chunk, word_count)
118
119     if X is None:     continue
120
121     # Get closest cluster indices
122     max_idx = helpers.sparse_matrix_to_cluster_indices(X, mu)
123
124     # Count cluster tags
125     for i, k in enumerate(max_idx):
126         for tag_idx in tags[i]:
127             cluster_tag_counts[k][tag_idx] += 1
128
129     # Assign tags to clusters
130     tags_labelled = []
131     cluster2tag = {}
132     for k, tag_counts in cluster_tag_counts.items():
133         tag_counts_sorted = sorted(tag_counts.items(), key=lambda
134             x: x[1], reverse=True)
135         for tag, count in tag_counts_sorted:
136             if tag not in tags_labelled:
137                 cluster2tag[k] = tag
138                 tags_labelled.append(tag)
139                 break
140
141     # Save cluster tags dict
142     config.data.save_cluster_tags(cluster_tags=cluster2tag)
143
144     # Save means
145     with open(config.paths.MU, 'wb') as f:
146         np.save(f, mu)
```

Include helper
functions and
config